

Sampling

Reservoir Sampling

Common Operations on Streams

- Sampling
- Filtering
- Counting Distinct
- Clustering

Sampling

- Extract samples from a stream such that they bear the approximately the same statistical properties- serves as the representative set.

Reservoir Sampling

Reservoir sampling is a family of randomized algorithms for randomly choosing a sample of k items from a list S containing n items, where n is either a very large or unknown number.

Assumptions-

n is large enough that the whole dataset does not fit into main memory, whereas k , the desired sample does.

Pseudo Code

```
array R[k]; // result
integer i, j;

// fill the reservoir array with the first k incoming elements
for each i in 1 to k do
    R[i] = S[i]
done;

// replace elements probabilistically
for each i in k+1 to length(S) do
    j = random(1, i); // important: inclusive range
    if j <= k then
        R[j] = S[i]
    end if;
Done;
```

Proof of Correctness

- Equal Likelihood:
 - Probability that any item $S[i]$ where $0 \leq i < n$ will be in final $R[]$ is k/n .

Proof of Correctness

- Let x_i be the i -th element and S_i be the solution obtained after examining the first i elements.
- RTS:-
 $\Pr[x_j \in S_i] = k/i$ for all $j \leq i$ with $k \leq i \leq n$.
This will imply that the probability that any element is in the final solution S_n is exactly k/n .

A Different Perspective

- The core insight behind reservoir sampling is that picking a random sample of size k is equivalent to *generating a random permutation (ordering) of the elements and picking the top k elements*.
- Associate a random float id with each element and pick the elements with the k largest ids. Since the ids induce a random ordering of the elements (assuming the ids are distinct), it is clear that the elements associated with the k largest ids form a random subset.

A Different Perspective

- The goal here is to incrementally keep track of the k elements with largest ids seen so far.
- (Streaming Sequential Setting)

```
import sys, random
from heapq import heappush, heapreplace

k = int(sys.argv[1])
H = []

for x in sys.stdin:
    r = random.random() # the randomly associated id.
    if len(H) < k: heappush(H, (r, x))
    elif r > H[0][0]: heapreplace(H, (r, x)) # replace

print ''.join([x for (r,x) in H]),
```

Map Reduce Approach

- Let K the number of samples you want. We'll assume that this is small enough to hold in memory on one of your nodes.
- Each mapper associates a random id with each element and keeps track of the top k elements.
- The top k elements of each mapper are then sent to a single reducer which will complete the job by extracting the top k elements among all.
- Data sent to the Reducer is now restricted to top k found in each mapper instead of the whole data set.

Mapper

```
# mapper.py
import sys, random
from heapq import heappush, heapreplace

k = int(sys.argv[1])
H = []

for x in sys.stdin:
    r = random.random() #randomly associated key, x is the value
    if len(H) < k: heappush(H, (r, x))
    elif r > H[0][0]: heapreplace(H, (r, x))

for (r, x) in H:
    #negating the keys, the reducer receives the elements from highest to lowest
    print '%f\t%s' % (-r, x),
```

Reducer

- Hadoop framework will automatically present the values to the reducer in order of keys from lowest to highest.

```
# reducer.py

import sys

k = int(sys.argv[1])
c = 0

for line in sys.stdin:
    (r, x) = line.split('\t', 1)
    print x, # emit the values
    c += 1
    if c == k: break
```

Compromise on linear time complexity?

- Each heap operation takes $O(\log k)$ time, so a trivial bound for the overall running time would be $O(n \log k)$.
- However, this bound can be improved as the heap replace operation is only executed when the i -th element is larger than the root of the heap.
- This happens only if the i -th element is one of the k largest elements among the first i elements, which happens with probability k/i .
- Therefore the expected number of heap replacements is $\sum_{i=k+1}^n k/i \approx k \log(n/k)$.
- The overall time complexity is then $O(n + k \log(n/k) \log k)$, which is substantially linear in n unless k is comparable to n .

Big Data it is!

- So far we worked under the assumption that the desired sample would fit into memory.
- After all, in the big data world, 1% of a huge dataset may still be too much to keep in memory!

Solution

- Use multiple reducers.
- The key idea is:
 - suppose we have ℓ buckets and generate a random ordering of the elements first by putting each element in a random bucket and then by generating a random ordering in each bucket.
 - The elements in the first bucket are considered smaller (with respect to the ordering) than the elements in the second bucket and so on.
 - if we want to pick a sample of size k , we can collect all of the elements in the first j buckets if they overall contain a number of elements t less than k , and then pick the remaining $k-t$ elements from the next bucket.
- Here ℓ is a parameter such that n/ℓ elements fit into memory.
- Note the key aspect that buckets can be processed distributively.

Mapper

- Mappers associate with each element an id (j,r) where j is a random index in $\{1,2,\dots,\ell\}$ to be used as key, and r is a random float for secondary sorting. In addition, mappers keep track of the number of elements with key less than j (for $1 \leq j \leq \ell$) and transmit this information to the reducers.

Mapper

```
# largeK_mapper.py

import sys, random
# number of buckets
l = int(sys.argv[1])
S = [0 for j in range(l)]

for x in sys.stdin:
    (j,r) = (random.randint(0,l-1), random.random()) #key
    S[j] += 1
    print '%d\t%f\t%s' % (j, r, x), #key, value pair

for j in range(l): # compute partial sums
    prev = 0 if j == 0 else S[j-1]
    S[j] += prev # number of elements with key less than j
    print '%d\t-1\t%d\t%d' % (j, prev, S[j]) # secondary key is -1 so reducer gets this first
```

Reducer

- The reducer associated with some key (bucket) j acts as follows: if the number of elements with key less or equal than j is less or equal than k then output all elements in bucket j ; otherwise, if the number of elements with key strictly less than j is $t < k$, then run a reservoir sampling to pick $k - t$ random elements from the bucket; in the remaining case, that is when the number of elements with key strictly less than j is at least k , don't output anything.

```
1
2 k = int(sys.argv[1])
3     line = sys.stdin.readline()
4 while line:
5     # Aggregate Mappers information
6     less_count, upto_count = 0, 0
7     (j, r, x) = line.split('\t', 2)
8     while float(r) == -1:
9         l, u = x.split('\t', 1)
10        less_count, upto_count = less_count + int(l), upto_count + int(u)
11        (j, r, x) = sys.stdin.readline().split('\t', 2)
12    n = upto_count - less_count # elements in bucket j
13
14    # Proceed with one of the three cases
15    if upto_count <= k: # in this case output the whole bucket
16        print x,
17        for i in range(n-1):
18            (j, r, x) = sys.stdin.readline().split('\t', 2)
19            print x,
20
21    elif less_count >= k: # in this case do not output anything
22        for i in range(n-1):
23            line = sys.stdin.readline()
24
25    else: # run reservoir sampling picking (k-less_count) elements
26        k = k - less_count
27        S = [x]
28        for i in range(1,n):
29            (j, r, x) = sys.stdin.readline().split('\t', 2)
30            if i < k:
31                S.append(x)
32            else:
33                r = random.randint(0,i-1)
34                if r < k: S[r] = x
35        print ''.join(S),
36    line = sys.stdin.readline()
```

Limitations

- Communication overhead of transferring *the whole* dataset from the mappers to the reducers as opposed to the k items only previously.
- The previous approach should be preferred if the sample size k fits in memory.